

Logistic Regression

CSC 461: Machine Learning

Fall 2020

Prof. Marco Alvarez
University of Rhode Island

Review of basic probability

So far ...

- Linear methods for classification
 - ✓ discrete outputs (perceptron)
- What if probabilistic outputs are needed?

Probability distributions



$P(W)$

W	P
sun	0.6
rain	0.1
fog	0.3
meteor	0.0

Random variable, domain, and probabilities

<https://inst.eecs.berkeley.edu/~cs188/fa19/assets/slides/lec13.pdf>

Joint distribution

- A distribution over a **set of random variables**
- Specifies probabilities for each **outcome**
- **Normalized:** sum to 1
- **Event** is a set of outcomes

P(T, W)		
T	W	P
hot	sun	0.4
hot	rain	0.1
cold	sun	0.2
cold	rain	0.3

$$\begin{array}{ll} P(\text{hot})? & .5 \\ P(\text{sun})? & .6 \\ P(\text{hot, rain})? & .1 \end{array}$$

<https://inst.eecs.berkeley.edu/~cs188/fa19/assets/slides/lec13.pdf>

Marginal distribution

P(X, Y)		
X	Y	P
+x	+y	0.2
+x	-y	0.3
-x	+y	0.4
-x	-y	0.1

$$\begin{aligned} P(x) &= \sum_y P(x, y) \\ P(y) &= \sum_x P(x, y) \end{aligned}$$

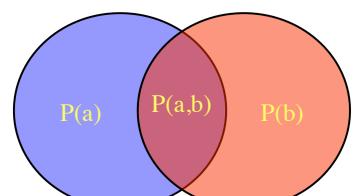
P(X)	
X	P
+x	
-x	

P(Y)	
Y	P
+y	
-y	

probability distribution of a subset of (marginal) random variables

<https://inst.eecs.berkeley.edu/~cs188/fa19/assets/slides/lec13.pdf>

Conditional probabilities



$$P(a|b) = \frac{P(a, b)}{P(b)}$$

$$P(X, Y)$$

X	Y	P
+x	+y	0.2
+x	-y	0.3
-x	+y	0.4
-x	-y	0.1

$$\begin{array}{lll} P(+x|+y)? & .2/.6 & P(-x|+y)? & .4/.6 & P(-y|+x)? & .3/.5 \end{array}$$

<https://inst.eecs.berkeley.edu/~cs188/fa19/assets/slides/lec13.pdf>

Conditional distribution

$$P(T, W)$$

T	W	P
hot	sun	0.4
hot	rain	0.1
cold	sun	0.2
cold	rain	0.3

SELECT the joint probabilities matching the evidence

T	W	P
cold	sun	0.2
cold	rain	0.3

NORMALIZE the selection (make it sum to one)

W	P
sun	0.4
rain	0.6

$$P(X, Y)$$

X	Y	P
+x	+y	0.2
+x	-y	0.3
-x	+y	0.4
-x	-y	0.1

SELECT the joint probabilities matching the evidence

NORMALIZE the selection (make it sum to one)

<https://inst.eecs.berkeley.edu/~cs188/fa19/assets/slides/lec13.pdf>

Logistic regression

- › Binary classification

- ✓ similar to the perceptron but uses a **logistic function** (type of sigmoid function, S-shaped)
 - ✓ models **probability** of output in terms of input

- › It is considered a **linear classifier**

- ✓ even though the ‘activation’ function is non-linear

- › It is a **discriminative model**

- ✓ models decision boundary directly, $P(y | \mathbf{x})$ in this case

Basics

$$\mathcal{D} = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})\}$$

$$\mathbf{x}^{(i)} \in \mathbb{R}^d, y^{(i)} \in \{-1, +1\}$$

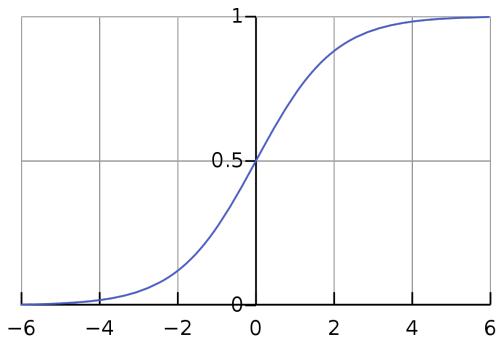
Logistic regression

Big picture

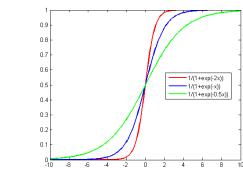
- › Define a hypothesis (classifier)
- › Define a loss function
- › Optimize with gradient descent
- › Predict the class with highest probability using the learned hypothesis

Logistic function

$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}$$



mapping \mathbb{R} to $[0,1]$
continuous and
differentiable



<https://alliance.seas.upenn.edu/~cis520/wiki/index.php?n=Lectures.Logistic>

Probabilistic interpretation

$$h_{\mathbf{w}}(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}}$$

(probability of
class +1)

$$P(y = +1 | \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}}$$

$$P(y = -1 | \mathbf{x}) = 1 - P(y = +1 | \mathbf{x})$$

(probability of
class -1)

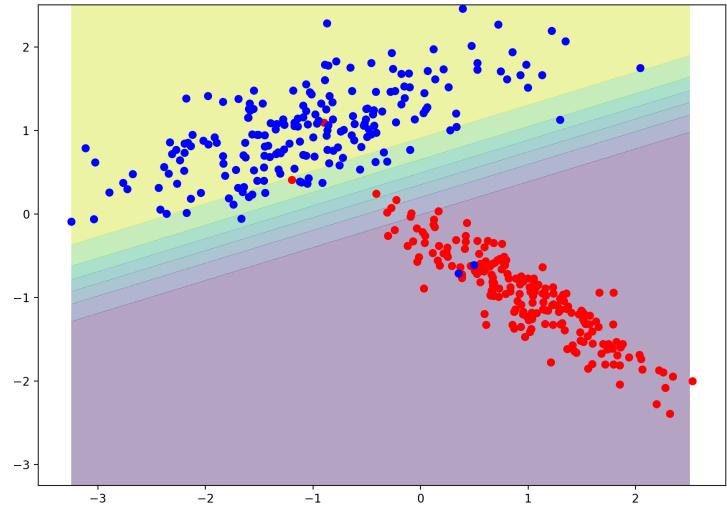
$$P(y = -1 | \mathbf{x}) = \sigma(-\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + e^{\mathbf{w}^T \mathbf{x}}}$$

Decision boundary

$$P(y = +1 | \mathbf{x}) = P(y = -1 | \mathbf{x}) = 0.5$$

Logistic regression has a linear decision boundary $\mathbf{w}^T \mathbf{x} = 0$

Linear decision boundary



Loss function

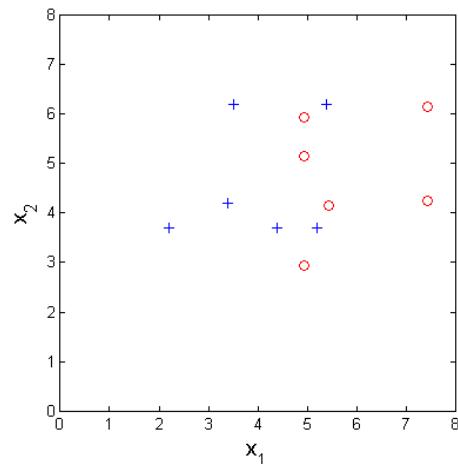
$$L(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \log \left(1 + e^{-y^{(i)} \mathbf{w}^T \mathbf{x}^{(i)}} \right)$$

Logistic regression
Cross-entropy loss

will be derived later ... when covering MLE

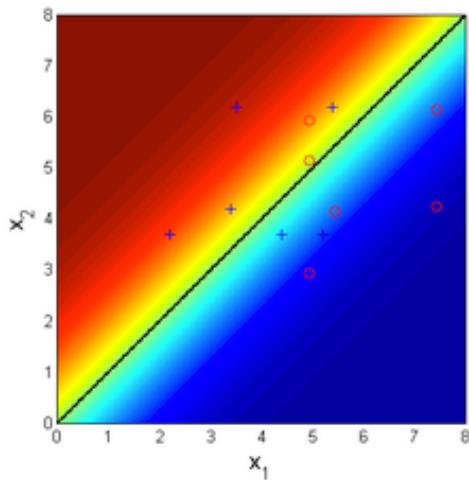
- ▶ no closed-form solution, but loss is convex
- ▶ can use gradient descent (or stochastic) or second-order methods

Example: 2d dataset



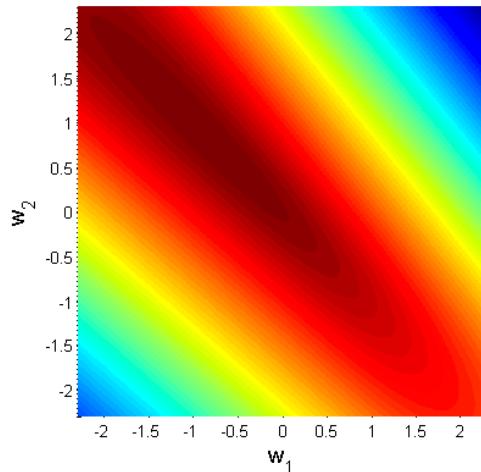
<https://alliance.seas.upenn.edu/~cis520/wiki/index.php?n=Lectures.Logistic>

Solution



<https://alliance.seas.upenn.edu/~cis520/wiki/index.php?n=Lectures.Logistic>

Example: loss function



plot shows contour lines in the space of parameters w_1 and w_2 ,
 w_0 is omitted

<https://alliance.seas.upenn.edu/~cis520/wiki/index.php?n=Lectures.Logistic>

Gradient

$$\nabla_{\mathbf{w}} L(\mathbf{w}) = \left[\frac{\partial L(\mathbf{w})}{\partial w_0}, \dots, \frac{\partial L(\mathbf{w})}{\partial w_d} \right]$$

$$\begin{aligned}\frac{\partial L(\mathbf{w})}{\partial w_j} &= \frac{\partial}{\partial w_j} \frac{1}{n} \sum_{i=1}^n \log \left(1 + e^{-y^{(i)} \mathbf{w}^T \mathbf{x}^{(i)}} \right) \\ &= -\frac{1}{n} \sum_{i=1}^n \sigma(-\mathbf{w}^T \mathbf{x}^{(i)}) y^{(i)} x_j^{(i)}\end{aligned}$$

How to classify new data?

- Once the final hypothesis $h_{\mathbf{w}}(\mathbf{x})$ is known ...
 - ✓ output the label of the most probable class
- Assign label **+1** to input instance \mathbf{x} , if $p(+1 | \mathbf{x}) > 0.5$ and label **-1** otherwise

Final remarks

- Simple classifier with probabilistic outputs
- Loss function is convex and can be trained with GD methods (no closed-form)
- Robust to overfitting
- Offers interpretability to weights (feature importance)
- However, decision boundary is still linear